

Integrity Report Review

Q&A-Driven Interactive Storytelling Using Large Language Models:
A Multimodal Framework for Personalised Narrative Generation

Anjali Madan Jha (124MTCM1008) · Guide: Dr. Vidyullata Devmane
Department of Computer Engineering, Shah & Anchor Kutchhi Engineering College

Analysed file: `final paper.pdf` (10 pp., pdfTeX-1.40.29, compiled 27 June 2026)

Review date: 18 August 2026

Summary

Two flags were raised on this paper. They have different answers, and it matters that they are not treated the same way.

Flag 1: Hidden text → **FALSE POSITIVE**. The 489 white characters are labels printed inside coloured boxes in three diagrams. Every one of the 60 white text spans sits on a dark filled rectangle and is plainly legible in the rendered PDF. Nothing is concealed. This flag can be contested with the evidence in Section 1.

Flag 2: AI writing → **PARTLY REAL**. Some of this is an artefact of the genre, because methods and results prose is uniform by nature and detectors score that heavily. But there is a genuine, repeated stylistic habit in the text, and it is worth fixing. Eight specific edits are given in Section 3.

Also found, unrelated to either flag: a factual contradiction. Section V.J reports ablation numbers that contradict Table V and Section V.F. This is more damaging than either integrity flag because an examiner comparing the two will catch it. See Section 4.

1. The hidden-text flag is a false positive

Where the white text is

All of it is confined to three figures on three consecutive pages. Pages 1–3 and 7–10 contain no white text at all.

Page	Figure	Spans	Chars
4	Fig. 2, Four-layer system architecture	28	293
5	Fig. 6, Image synthesis pipeline	14	77
6	Fig. 7, Multimodal audio signal flows	18	119
Total		60	489


The report cites 458 characters; the direct count is 489 across the same three pages. The small difference is a counting convention, most likely the em dashes and the `<img>` angle brackets.

Why it is not concealment


Each white span was tested against the vector fill shapes beneath it. **All 60 of 60 sit on a dark filled box**. Representative fill colours recovered from the page: RGB(43,43,43) charcoal,

RGB(86, 0, 169) purple, RGB(103, 26, 178) violet, RGB(64, 64, 191) indigo. White is the only readable text colour on those fills.

The figures below are rendered directly from the flagged pages at 4× magnification. The text the report calls hidden is the text you can read here.



Page 4, Fig. 2. All 28 flagged spans on this page are these box labels.



Page 6, Fig. 7. All 18 flagged spans on this page are these box labels.

Checks for genuine concealment: all negative

The four techniques actually used to hide text in a PDF were each tested across all ten pages.

Concealment technique	Instances found
Invisible text render mode (3 Tr / 7 Tr operators)	0
Sub-legible font size (< 4 pt)	0
Text positioned outside the visible page rectangle	0
Text covered by an opaque image	0

Complete colour census

Every text character in the document, grouped by colour. There is no fourth category and no white body paragraph.

Colour	Chars	What it is
RGB(0,0,0)	42,669	Body text, captions, tables
RGB(255,255,255)	489	Diagram box labels (Figs. 2, 6, 7)
RGB(133,133,133)	26	Figure annotation (“Modular Core”)
RGB(115,115,115)	8	Figure annotation

There is no keyword stuffing, no inflated word count, and no text that is invisible to a reader but visible to a similarity checker.

Recommended response

Contest the flag. One sentence is enough: *the white characters are labels inside coloured boxes in Figures 2, 6 and 7, and they are visible in the rendered PDF*. Point the reviewer at those three figures. Nothing in the paper needs to change.

If the committee requires the flag cleared regardless, compile each of the three diagrams as a standalone PDF and convert its text to vector outlines (`gs -o fig-flat.pdf -dNoOutputFonts -sDEVICE=pdfwrite fig.pdf`), then include the flattened file. This removes the text layer. It is worth noting this makes the figures slightly worse for screen readers in order to satisfy a detector that has misread them.

2. The AI-writing flag

What is genre artefact

The flagged region is Sections IV.E through V.F: endpoint names, table pointers, percentages. This prose is uniform and low-entropy by nature, which is what AI detectors measure. Constructions such as “Three evaluation axes were used” or “Table IV reports mean latencies” are correct IEEE register. Passive voice and formulaic figure pointers are how methods sections are written, and no rewrite makes them read as less mechanical to a classifier. These tools have well-documented false-positive rates on exactly this genre.

What is genuinely worth fixing

The following patterns are real and repeat throughout the paper, not only in the flagged region.

Pattern	Count	Assessment
Participial clause-tails (<i>, confirming...</i>)	13	Dominant tell
Em dashes	20	Secondary
Evaluative filler in results sections	4	Editorialising
Caption text duplicating body text	Figs. 9, 10	Also a reviewer complaint

The strongest signal is clustering: Sections V.D, V.E and V.F each close on a participial tail (*confirming...*, *indicating...*, *confirming...*). Three in a row across consecutive subsections is what reads as machine-generated. Removing them is the single highest-value change.

3. Eight edits

Facts and figures are unchanged. Register stays formal, because plain and neutral is the correct human voice for an IEEE paper. These apply directly in the Overleaf editor.

1. Section V.D

```
FIND  , confirming Q&A elicitation as the core value driver
REPL  . Both derive directly from the Q&A step
```

2. Section V.E

FIND , indicating a need for improved discoverability cues

REPL . Both sit behind a secondary tab, which is the most likely explanation

3. Section V.F

FIND , confirming that deeper Q&A elicitation adds measurable value

REPL (delete it; the numbers already make the point)

4. Section IV.E

FIND (enabling session replay)

REPL , which is what makes session replay possible

FIND (persisting provider and temperature settings)

REPL , which holds the provider and temperature the user selected

5. Section V.D

FIND disaggregates satisfaction

REPL breaks satisfaction down

6. Section V.B (removes the em-dash pair). In the .tex the em dash is typed as three hyphens, shown below as ---.

FIND 1.2 s continuation---within the 5 s design target---with continuation faster

REPL 1.2 s continuation, both inside the 5 s design target. Continuations are faster

7. Section V.J · highest priority (see Section 4)

FIND improves personalisation by ≈ 0.95 Likert points while coherence stays flat

REPL improves personalisation by ≈ 1.4 Likert points, with coherence rising more modestly (0.7 points)

8. Captions of Figs. 9 and 10. Rewrite so they describe what the axes show rather than restating their paragraphs. At present the same sentence appears twice on one page. For example:

REPL Fig. 10. Per-feature adoption across 15 user-study sessions. Each bar is the share of sessions in which the feature was used at least once.

4. Priority: Section V.J contradicts Table V

This is independent of both integrity flags and should be fixed first.

Table V (question-budget ablation) gives these values:

Budget	Coherence	Creativity	Personalisation	Time
4 (Quick)	3.6	3.8	3.2	45 s
8 (Balanced)	3.9	4.1	3.8	1.5 m
12 (Detailed)	4.2	4.2	4.3	2.5 m
17 (Full)	4.3	4.1	4.6	4.6 m

Section V.F reports this correctly: personalisation rises ≈ 1.4 points ($3.2 \rightarrow 4.6$) and coherence rises $3.6 \rightarrow 4.3$.

Section V.J then states: “improves personalisation by ≈ 0.95 Likert points while coherence stays flat.”

Both halves are wrong against the paper’s own table. The personalisation gain is 1.4, not 0.95. Coherence rises 0.7 points, which is not flat. Correct V.J as given in edit 7.

Secondary observation. Creativity falls from 4.2 to 4.1 between the 12-question and 17-question budgets. The text does not mention this, but it supports the existing recommendation of 12 questions as the default and is worth a clause.

5. Status of the source files

The `.tex` source for the reviewed paper is not on this machine. It exists here only as a compiled PDF. Two older copies of the paper source were located:

Check	projects pack copy	Anjali maam/latex	Flagged PDF
Lines	898	909	—
Title	“Q&A- <i>Based...</i> ”	“Q&A- <i>Based...</i> ”	“Q&A- <i>Driven...</i> ”
<code>tikzpicture</code>	0	0	3
<code>core value driver</code>	0	0	present
0.95 Likert	0	0	present
CoNoder / Salmaze refs	0	0	cited

Searches covered every `.tex` and `.bib` in the home directory, plus the interior of every `.docx`, `.pptx` and `.zip` on the Desktop. None contain the flagged sentences or references [32]–[42].

The reviewed version therefore adds a rewritten related-work section, three TikZ diagrams and a changed title relative to both local copies. To apply these edits to the real paper, export the source from Overleaf (Menu → Download → Source).

Caution. Do not rebuild from the local copies. Doing so loses the newer related-work section and the TikZ figures, and would submit a paper whose title differs from the one that was reviewed.

6. A note on disclosure

Whether any part of this text was drafted with AI assistance cannot be determined from the file, and is not assumed here. If it was, editing the stylistic markers does not substitute for whatever disclosure the institution requires, and following that policy is the safer course than reducing a detector score. If it was not, the eight edits above are ordinary copy-editing that improve the paper, and the caption duplication and the Table V contradiction were worth catching either way.

Evidence for Section 1 was produced with PyMuPDF 1.26.5 over `final paper.pdf`: per-span colour extraction, vector fill-shape intersection at each span centre, content-stream scan for text render-mode operators, and $4\times$ region rendering. Figure crops are unmodified renders of pages 4 and 6.